

RL Adventure

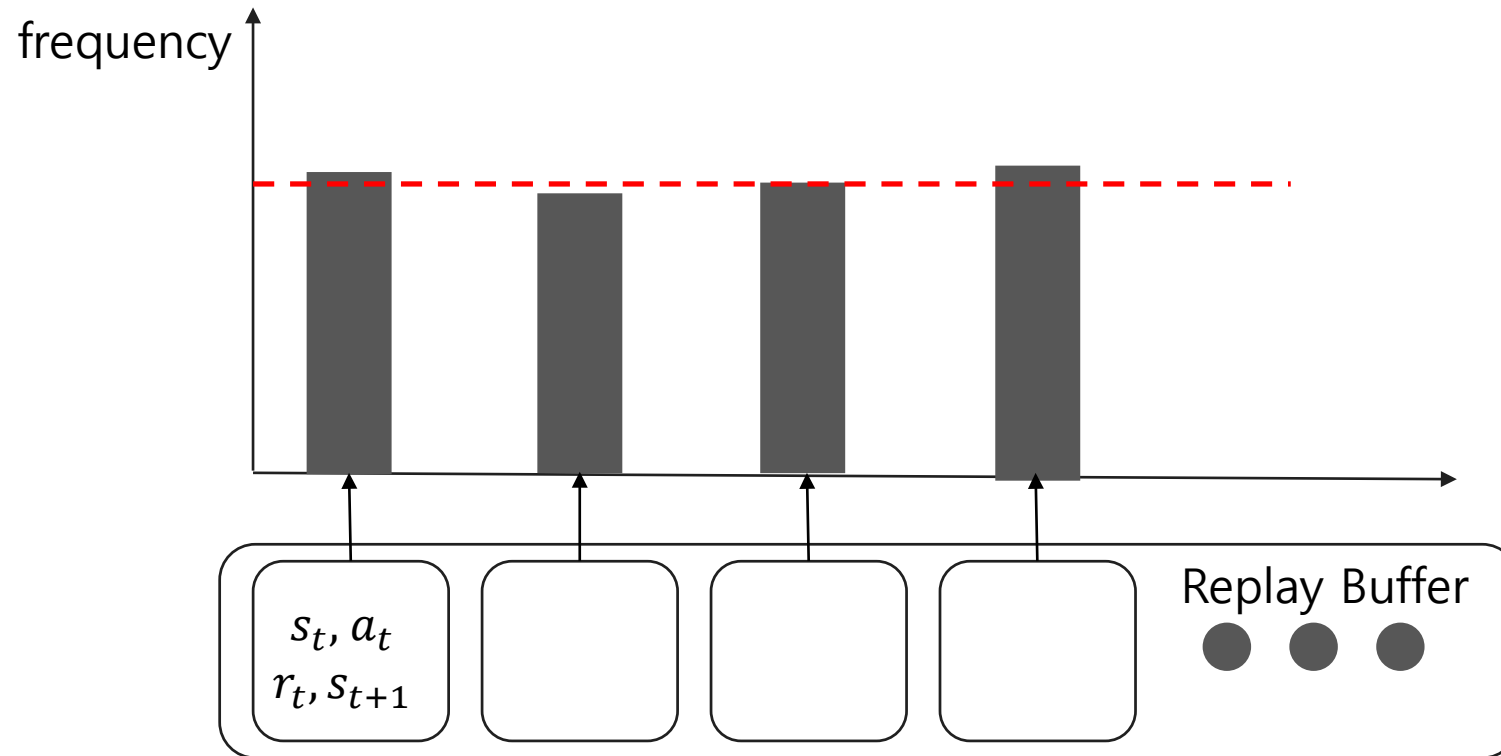
PER and NoisyNet

양홍선


PER

Prioritized Experience Replay


Replay Memory



Replay Memory



Replay Memory



Design of Replay Memory

Which experiences to store


Which experiences to replay

Design of Replay Memory

Which experiences to store

Which experiences to replay

A Motivating Example



Two actions: 'right($\xrightarrow{\text{green}} \rightarrow$)' and 'wrong($\xrightarrow{\text{red}} \rightarrow$)'

The environment requires an exponential number of random steps until the first non-zero reward

The most relevant transitions are hidden in a mass of highly redundant failure cases

How?

Prioritizing with TD-Error

A transition's TD error δ

→ how 'surprising' or unexpected the transition is

Weakness

A low TD-Error on first visit may not be replayed for a long time

The PER with TD-Error is sensitive to noise spikes

Greedy prioritization focuses on a small subset of the experience

Stochastic Sampling!

Stochastic Prioritization

Proportional prioritization

- $p_i = |\delta_i| + \epsilon$
- $P(i) = \frac{p_i^\alpha}{\sum_k p_k^\alpha}$
 - $p_i > 0$: the priority of transition i
 - α : determines how much prioritization is used
- Sum-tree

Stochastic Prioritization

Rank-based prioritization

- $p_i = \frac{1}{\text{rank}(i)}$
 - $\text{rank}(i)$ is the rank of transition i when the replay memory is sorted according to $|\delta_i|$
 - More robust
- Binary heap

Annealing the Bias

- Importance-Sampling (IS) weights

- $w_i = \left(\frac{1}{N} \frac{1}{P(i)} \right)^\beta$

- Normalize: $\frac{1}{\max_i w_i}$

- $\Delta \leftarrow \Delta + \boxed{w_i \cdot \delta_i} \cdot \nabla_\theta Q(S_{i-1}, A_{i-1})$

Proportional prioritization (without sum-tree)

```
class NaivePrioritizedBuffer(object):  
    def __init__(self, capacity, prob_alpha=0.6):  
  
    def push(self, state, action, reward, next_state, done):  
  
    def sample(self, batch_size, beta=0.4):  
  
    def update_priorities(self, batch_indices, batch_priorities):  
  
    def __len__(self):
```



```
class NaivePrioritizedBuffer(object):
    def __init__(self, capacity, prob_alpha):

    def push(self, state, action, reward,

    def sample(self, batch_size, beta=0.4

    def update_priorities(self, batch_ind

    def __len__(self):
```

```
def __init__(self, capacity, prob_alpha=0.6):
    self.prob_alpha = prob_alpha
    self.capacity = capacity
    self.buffer = []
    self.pos = 0
    self.priorities = np.zeros((capacity,), dtype=np.float32)
```

```

class NaivePrioritizedBuffer(object):
    def __init__(self, capacity, prob_alpha):

    def push(self, state, action, reward,

    def sample(self, batch_size, beta=0.4

    def update_priorities(self, batch_ind

    def __len__(self):

```

$$P(i) = \frac{p_i^\alpha}{\sum_k p_k^\alpha}$$

```

def __init__(self, capacity, prob_alpha=0.6):
    self.prob_alpha = prob_alpha
    self.capacity = capacity
    self.buffer = []
    self.pos = 0
    self.priorities = np.zeros((capacity,), dtype=np.float32)

```

```

class NaivePrioritizedBuffer(object):
    def __init__(self, capacity, probab_als):

    def push(self, state, action, reward, next_state, done):

    def sample(self, batch_size, beta=0.4):

    def update_priorities(self, batch_indices, priorities):

    def __len__(self):

```

```

def push(self, state, action, reward, next_state, done):
    assert state.ndim == next_state.ndim
    state = np.expand_dims(state, 0)
    next_state = np.expand_dims(next_state, 0)

    max_prio = self.priorities.max() if self.buffer else 1.0

    if len(self.buffer) < self.capacity:
        self.buffer.append((state, action, reward, next_state, done))
    else:
        self.buffer[self.pos] = (state, action, reward, next_state, done)

    self.priorities[self.pos] = max_prio
    self.pos = (self.pos + 1) % self.capacity

```

```

class NaivePrioritizedBuffer(object):
    def __init__(self, capacity, probab_ald

    def push(self, state, action, reward,

    def sample(self, batch_size, beta=0.4

    def update_priorities(self, batch_ind

    def __len__(self):

```

```

def push(self, state, action, reward, next_state, done):
    assert state.ndim == next_state.ndim
    state = np.expand_dims(state, 0)
    next_state = np.expand_dims(next_state, 0)

```

최소 한번은 replay

```

max_prio = self.priorities.max() if self.buffer else 1.0

```

```

if len(self.buffer) < self.capacity:
    self.buffer.append((state, action, reward, next_state, done))
else:
    self.buffer[self.pos] = (state, action, reward, next_state, done)

self.priorities[self.pos] = max_prio
self.pos = (self.pos + 1) % self.capacity

```

```

class NaivePrioritizedBuffer(object):
    def __init__(self, capacity, prob_alpha):

    def push(self, state, action, reward, done):

    def sample(self, batch_size, beta=0.4):

    def update_priorities(self, batch_indices, weights):

    def __len__(self):

```

```

def sample(self, batch_size, beta=0.4):
    if len(self.buffer) == self.capacity:
        prios = self.priorities
    else:
        prios = self.priorities[:self.pos]

    probs = prios ** self.prob_alpha
    probs /= probs.sum()

    indices = np.random.choice(len(self.buffer), batch_size, p=probs)
    samples = [self.buffer[idx] for idx in indices]

    total = len(self.buffer)
    weights = (total * probs[indices]) ** (-beta)
    weights /= weights.max()
    weights = np.array(weights, dtype=np.float32)

    batch = list(zip(*samples))
    states = np.concatenate(batch[0])
    actions = batch[1]
    rewards = batch[2]
    next_states = np.concatenate(batch[3])
    dones = batch[4]

    return states, actions, rewards, next_states, dones, indices, weights

```

```

class NaivePrioritizedBuffer(object):
    def __init__(self, capacity, prob_alpha):

    def push(self, state, action, reward,

    def sample(self, batch_size, beta=0.4):

    def update_priorities(self, batch_indices):

    def __len__(self):

```

```

def sample(self, batch_size, beta=0.4):
    if len(self.buffer) == self.capacity:
        prios = self.priorities
    else:
        prios = self.priorities[:self.pos]

```

$$P(i) = \frac{p_i^\alpha}{\sum_k p_k^\alpha}$$

```

probs = prios ** self.prob_alpha
probs /= probs.sum()

```

```

indices = np.random.choice(len(self.buffer), batch_size, p=probs)
samples = [self.buffer[idx] for idx in indices]

```

```

total = len(self.buffer)
weights = (total * probs[indices]) ** (-beta)
weights /= weights.max()
weights = np.array(weights, dtype=np.float32)

```

```

batch = list(zip(*samples))
states = np.concatenate(batch[0])
actions = batch[1]
rewards = batch[2]
next_states = np.concatenate(batch[3])
dones = batch[4]

```

```

return states, actions, rewards, next_states, dones, indices, weights

```

```

class NaivePrioritizedBuffer(object):
    def __init__(self, capacity, prob_alpha):

    def push(self, state, action, reward,

    def sample(self, batch_size, beta=0.4):

    def update_priorities(self, batch_indices):

    def __len__(self):

```

```

def sample(self, batch_size, beta=0.4):
    if len(self.buffer) == self.capacity:
        prios = self.priorities
    else:
        prios = self.priorities[:self.pos]

    probs = prios ** self.prob_alpha
    probs /= probs.sum()

    indices = np.random.choice(len(self.buffer), batch_size, p=probs)
    samples = [self.buffer[idx] for idx in indices]

```

IS weights

```

total = len(self.buffer)
weights = (total * probs[indices]) ** (-beta)
weights /= weights.max()
weights = np.array(weights, dtype=np.float32)

```

```

batch = list(zip(*samples))
states = np.concatenate(batch[0])
actions = batch[1]
rewards = batch[2]
next_states = np.concatenate(batch[3])
dones = batch[4]

```

```

return states, actions, rewards, next_states, dones, indices, weights

```

```
class NaivePrioritizedBuffer(object):
    def __init__(self, capacity, probab_alp

    def push(self, state, action, reward,

    def sample(self, batch_size, beta=0.4

    def update_priorities(self, batch_ind

    def __len__(self):
```

```
def update_priorities(self, batch_indices, batch_priorities):
    for idx, prio in zip(batch_indices, batch_priorities):
        self.priorities[idx] = prio

def __len__(self):
    return len(self.buffer)
```



```

class NaivePrioritizedBuffer(object):
    def __init__(self, capacity, probab_als):
        self.capacity = capacity
        self.buffer = []
        self.priorities = []

    def push(self, state, action, reward, next_state):
        self.buffer.append((state, action, reward, next_state))
        self.priorities.append(1)

    def sample(self, batch_size, beta=0.4):
        batch_indices = []
        batch_priorities = []
        for _ in range(batch_size):
            idx = np.random.choice(len(self.buffer))
            batch_indices.append(idx)
            batch_priorities.append(self.priorities[idx])

    def update_priorities(self, batch_indices, batch_priorities):
        for idx, prio in zip(batch_indices, batch_priorities):
            self.priorities[idx] = prio

    def __len__(self):
        return len(self.buffer)

```

```

def update_priorities(self, batch_indices, batch_priorities): TD-Error로 업데이트
    for idx, prio in zip(batch_indices, batch_priorities):
        self.priorities[idx] = prio

def __len__(self):
    return len(self.buffer)

```



NoisyNet

Noisy Networks for Exploration

Exploitation

Exploration

Exploitation



Efficient Exploration

Exploration methods

ϵ -greedy

일정 확률 (ϵ) 만큼 무작위로 행동


Entropy regularization

Loss에 추가하는 패널티로 한쪽으로 치우치지 않게 함

$$-\sum_a \pi(s,a) \log \pi(s,a)$$


ϵ – greedy, Entropy regularization

ϵ – greedy, Entropy regularization




Random perturbations

ϵ – greedy, Entropy regularization



Random perturbations



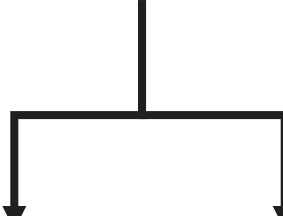
Hard to the large-scale behavioural patterns

NoisyNet!!

NoisyNet learn perturbations of the network weights are used to drive exploration

$$\theta := \mu + \Sigma \odot \epsilon$$

Learnable parameters



The diagram shows a vertical line descending from the text 'Learnable parameters'. This line splits into two horizontal branches, each ending in a downward-pointing arrow. The left arrow points to the Greek letter μ in the equation $\theta := \mu + \Sigma \odot \epsilon$. The right arrow points to the \odot symbol in the same equation.

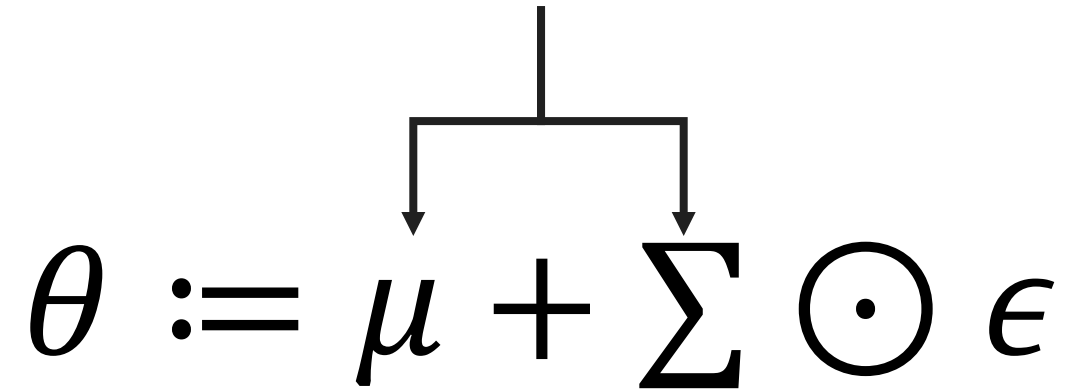
$$\theta := \mu + \Sigma \odot \epsilon$$



The diagram shows a vertical line with an upward-pointing arrow at the bottom, connecting the text 'Noise variables' to the ϵ term in the equation above.

Noise variables

$\zeta := (\mu, \Sigma)$ Learnable parameters




The diagram shows a hierarchical structure. A vertical line descends from the text ' $\zeta := (\mu, \Sigma)$ Learnable parameters'. This line splits into two horizontal branches. From each branch, an arrow points down to the terms μ and Σ in the equation $\theta := \mu + \Sigma \odot \epsilon$. The symbol \odot is a circle with a dot inside.

$$\theta := \mu + \Sigma \odot \epsilon$$

Noise variables




$$y = wx + b$$



$$y := (\mu^w + \sigma^w \odot \epsilon^w)x + \mu^b + \sigma^b \odot \epsilon^b$$

NoisyNet

- p inputs and q outputs
- Independent Gaussian noise
 - Using an independent Gaussian noise entry per weight
 - $pq+q$
- Factorised Gaussian noise
 - Using an independent noise per each output and input
 - $p+q$



Loss

$$L(\theta) = \mathbb{E} \left[\mathbb{E}_{(x,a,r,y) \sim D} [r + \gamma \max_{b \in A} Q(y, b; \theta^-) - Q(x, a; \theta)]^2 \right]$$

$$\bar{L}(\zeta) = \mathbb{E} \left[\mathbb{E}_{(x,a,r,y) \sim D} [r + \gamma \max_{b \in A} Q(y, b, \epsilon'; \zeta^-) - Q(x, a, \epsilon; \zeta)]^2 \right]$$

Loss

$$L(\theta) = \mathbb{E} \left[\mathbb{E}_{(x,a,r,y) \sim D} [r + \gamma \max_{b \in A} Q(y, b; \theta^-) - Q(x, a; \theta)]^2 \right]$$


$$\bar{L}(\zeta) = \mathbb{E} \left[\mathbb{E}_{(x,a,r,y) \sim D} [r + \gamma \max_{b \in A} Q(y, b; \epsilon'; \zeta^-) - Q(x, a; \epsilon; \zeta)]^2 \right]$$


Initialisation of NoisyNet

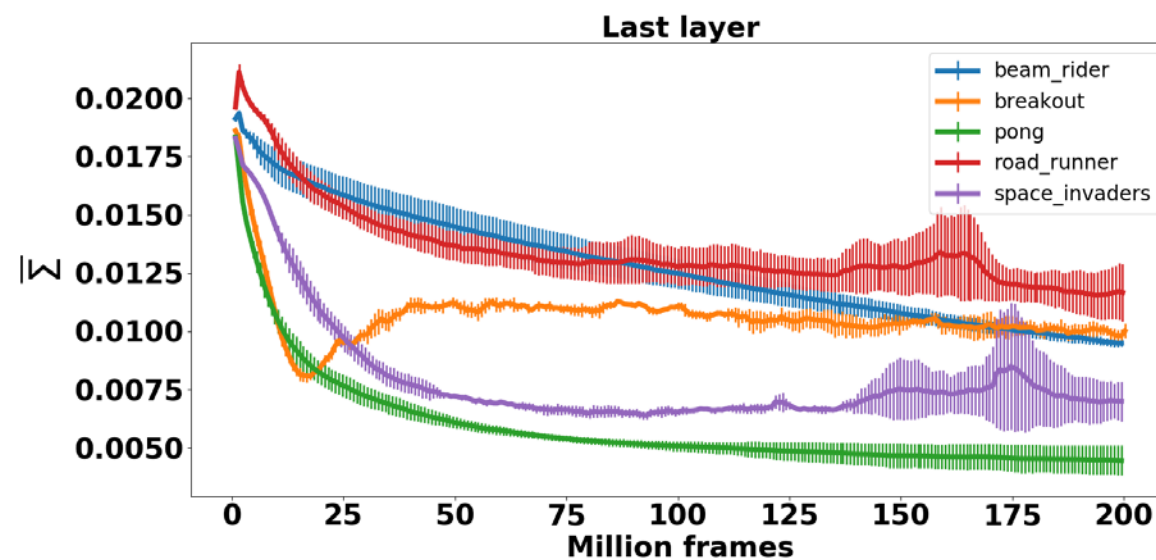

- An unfactorized NoisyNet

- $\mu_{i,j} \sim u \left[-\sqrt{\frac{3}{p}}, +\sqrt{\frac{3}{p}} \right]$
 - p : The number of inputs
- $\sigma_{i,j} = 0.017$


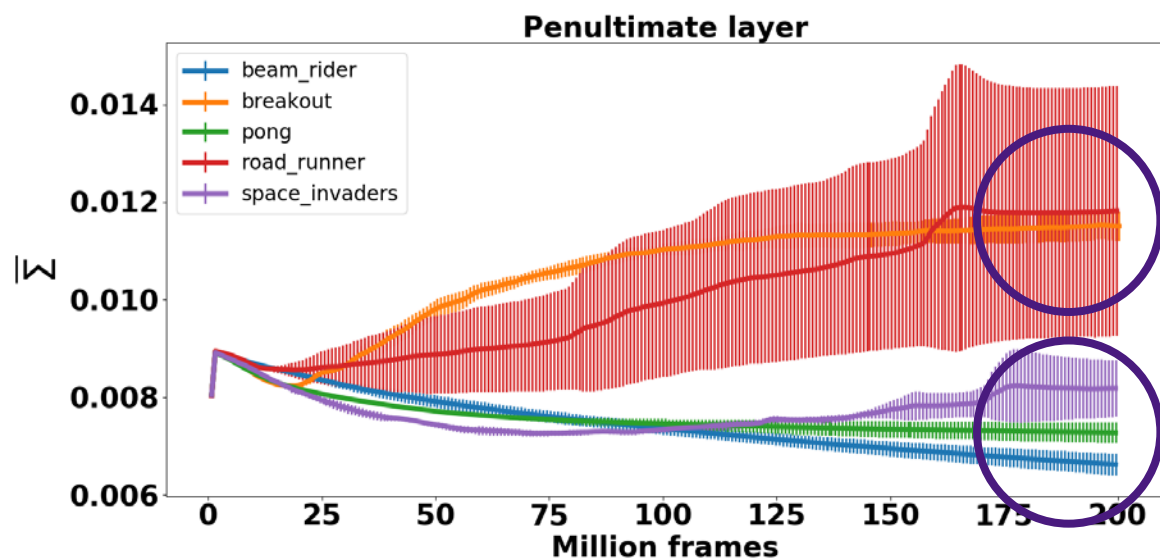
- Factorised NoisyNet

- $\mu_{i,j} \sim u \left[-\frac{1}{\sqrt{p}}, +\frac{1}{\sqrt{p}} \right]$
- $\sigma_{i,j} = \frac{\sigma_0}{\sqrt{p}}$
 - $\sigma_0 = 0.5$

The learning curves of the average noise parameter $\bar{\Sigma}$



The learning curves of the average noise parameter $\bar{\Sigma}$



Factorised NoisyNet

```
class NoisyLinear(nn.Module):  
    def __init__(self, in_features, out_features, std_init=0.4):  
  
    def forward(self, x):  
  
    def reset_parameters(self):  
  
    def reset_noise(self):  
  
    def _scale_noise(self, size):
```

```
class NoisyLinear(nn.Module):
    def __init__(self, in_features

    def forward(self, x):

    def reset_parameters(self):

    def reset_noise(self):

    def _scale_noise(self, size):
```

```
def __init__(self, in_features, out_features, std_init=0.5):
    super(NoisyLinear, self).__init__()

    self.in_features = in_features
    self.out_features = out_features
    self.std_init = std_init

    self.weight_mu = nn.Parameter(torch.FloatTensor(out_features, in_features))
    self.weight_sigma = nn.Parameter(torch.FloatTensor(out_features, in_features))
    self.register_buffer('weight_epsilon', torch.FloatTensor(out_features, in_features))

    self.bias_mu = nn.Parameter(torch.FloatTensor(out_features))
    self.bias_sigma = nn.Parameter(torch.FloatTensor(out_features))
    self.register_buffer('bias_epsilon', torch.FloatTensor(out_features))

    self.reset_parameters()
    self.reset_noise()
```

```

class NoisyLinear(nn.Module):
    def __init__(self, in_features

    def forward(self, x):

    def reset_parameters(self):

    def reset_noise(self):

    def _scale_noise(self, size):

```

```

def __init__(self, in_features, out_features, std_init=0.5):
    super(NoisyLinear, self).__init__()

    self.in_features = in_features
    self.out_features = out_features
    self.std_init = std_init

```

Learnable parameters

```

self.weight_mu = nn.Parameter(torch.FloatTensor(out_features, in_features))
self.weight_sigma = nn.Parameter(torch.FloatTensor(out_features, in_features))
self.register_buffer('weight_epsilon', torch.FloatTensor(out_features, in_features))

```

```

self.bias_mu = nn.Parameter(torch.FloatTensor(out_features))
self.bias_sigma = nn.Parameter(torch.FloatTensor(out_features))
self.register_buffer('bias_epsilon', torch.FloatTensor(out_features))

```

```

self.reset_parameters()
self.reset_noise()

```



```

class NoisyLinear(nn.Module):
    def __init__(self, in_features,

    def forward(self, x):

    def reset_parameters(self):

    def reset_noise(self):

    def _scale_noise(self, size):

```

Factorised NoisyNet

$$\mu_{i,j} \sim u \left[-\frac{1}{\sqrt{p}}, +\frac{1}{\sqrt{p}} \right]$$

```

def reset_parameters(self):
    mu_range = 1 / math.sqrt(self.weight_mu.size(1))

    self.weight_mu.data.uniform_(-mu_range, mu_range)
    self.weight_sigma.data.fill_(self.std_init / math.sqrt(self.weight_sigma.size(1)))

    self.bias_mu.data.uniform_(-mu_range, mu_range)
    self.bias_sigma.data.fill_(self.std_init / math.sqrt(self.bias_sigma.size(0)))

def reset_noise(self):
    epsilon_in = self._scale_noise(self.in_features)
    epsilon_out = self._scale_noise(self.out_features)

    self.weight_epsilon.copy_(epsilon_out.ger(epsilon_in))
    self.bias_epsilon.copy_(self._scale_noise(self.out_features))

def _scale_noise(self, size):
    x = torch.randn(size)
    x = x.sign().mul(x.abs().sqrt())
    return x

```

```

class NoisyLinear(nn.Module):
    def __init__(self, in_features,

    def forward(self, x):

    def reset_parameters(self):

    def reset_noise(self):

    def _scale_noise(self, size):

```

```

def reset_parameters(self):
    mu_range = 1 / math.sqrt(self.weight_mu.size(1))

    self.weight_mu.data.uniform_(-mu_range, mu_range)
    self.weight_sigma.data.fill_(self.std_init / math.sqrt(self.weight_sigma.size(1)))

    self.bias_mu.data.uniform_(-mu_range, mu_range)
    self.bias_sigma.data.fill_(self.std_init / math.sqrt(self.bias_sigma.size(0)))

def reset_noise(self):
    epsilon_in = self._scale_noise(self.in_features)
    epsilon_out = self._scale_noise(self.out_features)

    self.weight_epsilon.copy_(epsilon_out.ger(epsilon_in))
    self.bias_epsilon.copy_(self._scale_noise(self.out_features))

def _scale_noise(self, size):
    x = torch.randn(size)
    x = x.sign().mul(x.abs().sqrt())
    return x

```

Factorised

```
class NoisyLinear(nn.Module):
    def __init__(self, in_features,

    def forward(self, x):

    def reset_parameters(self):

    def reset_noise(self):

    def _scale_noise(self, size):
```

```
def forward(self, x):
    if self.training:
        weight = self.weight_mu + self.weight_sigma.mul(self.weight_epsilon)
        bias = self.bias_mu + self.bias_sigma.mul(self.bias_epsilon)
    else:
        weight = self.weight_mu
        bias = self.bias_mu

    return F.linear(x, weight, bias)
```

```

class NoisyLinear(nn.Module):
    def __init__(self, in_features,

    def forward(self, x):

    def reset_parameters(self):

    def reset_noise(self):

    def _scale_noise(self, size):

```

```

def forward(self, x):
    if self.training:

```

$$\theta := \mu + \sum \odot \epsilon$$

```

        weight = self.weight_mu + self.weight_sigma.mul(self.weight_epsilon)
        bias    = self.bias_mu    + self.bias_sigma.mul(self.bias_epsilon)

```

```

    else:

```

```

        weight = self.weight_mu
        bias    = self.bias_mu

```

```

    return F.linear(x, weight, bias)

```

Code: <https://github.com/higgsfield/RL-Adventure>

PER: <https://arxiv.org/abs/1511.05952>

NoisyNet: <https://arxiv.org/abs/1706.10295>

감사합니다

Q&A